

DataStax Enterprise

DSE Core Architecture & Modelling

Negib Marhoul
Solutions Engineer
negib.marhoul@datastax.com

Agenda

1	Topology and Architecture
2	Tuneable Consistency
3	Query Based Modelling

An abstract network diagram consisting of numerous small blue dots (nodes) connected by thin, light blue lines. The nodes are scattered across the right side of the image, with a higher density in the lower half. Some lines are thicker than others, and the overall structure suggests a complex, interconnected system.

Design for failure and nothing will fail

Apache Cassandra™ Architecture

Cluster layer

- Amazon DynamoDB paper
- masterless architecture


Data-store layer

- Google Big Table paper
- Columns / columns family



Topology


Cluster



Master-less Always-On, Scalable, Distributed

Continues Availability


- No master, Master Less
- Topology discovery
- Client topology awareness



AlwaysOn

Linear Scalability


- Scale out
- tunable consistency
- Runs on Commodity hardware



Linear Scalability


Built-in data distribution

- Shared-Nothing Architecture
- coordination free
- Automatic data distribution



On Premise, Cloud or Hybrid



Scale-Up Linearity



<http://techblog.netflix.com/2011/11/benchmarking-cassandra-scalability-on.html>

Distributed Cluster Or Hub'n Spoke Architecture

Multi-DC DSE Cluster *Central Cluster*

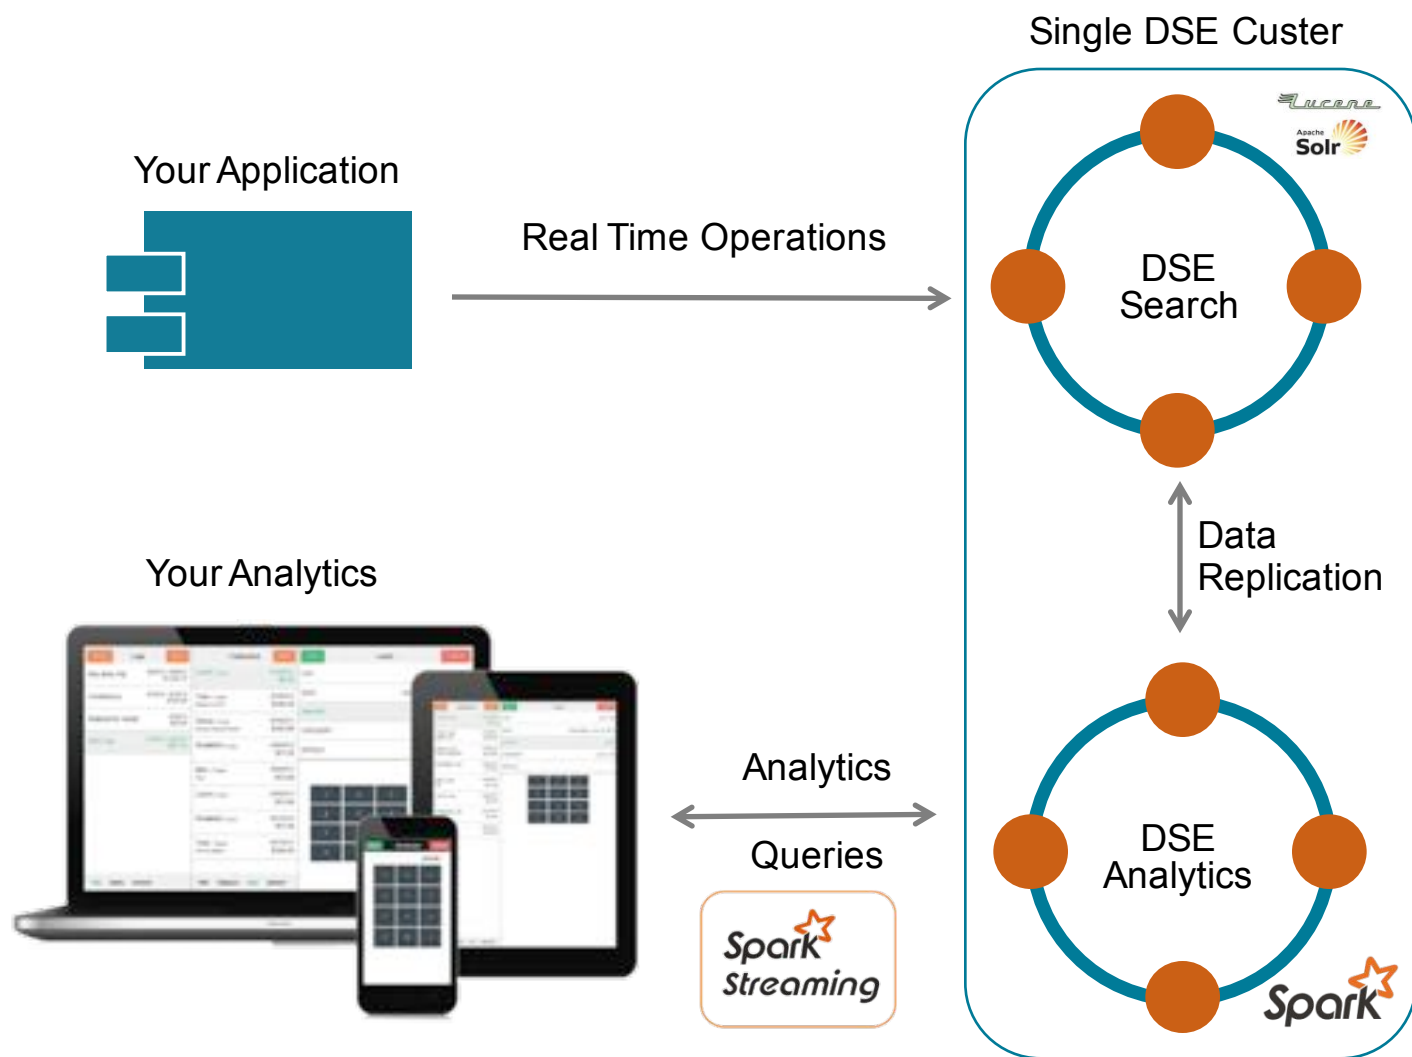


What is Advanced Replication

Advanced Replication supports:

- Many edge clusters replicating to a central hub
- Central hub replicating out to many edge clusters
- Consistent or sporadic connectivity – “store and forward”
- Active queries at the edge, as well as replicating to the hub
- Filter rules about which data is forwarded
- Prioritized streams for limited bandwidth situations

DSE Reference Architecture



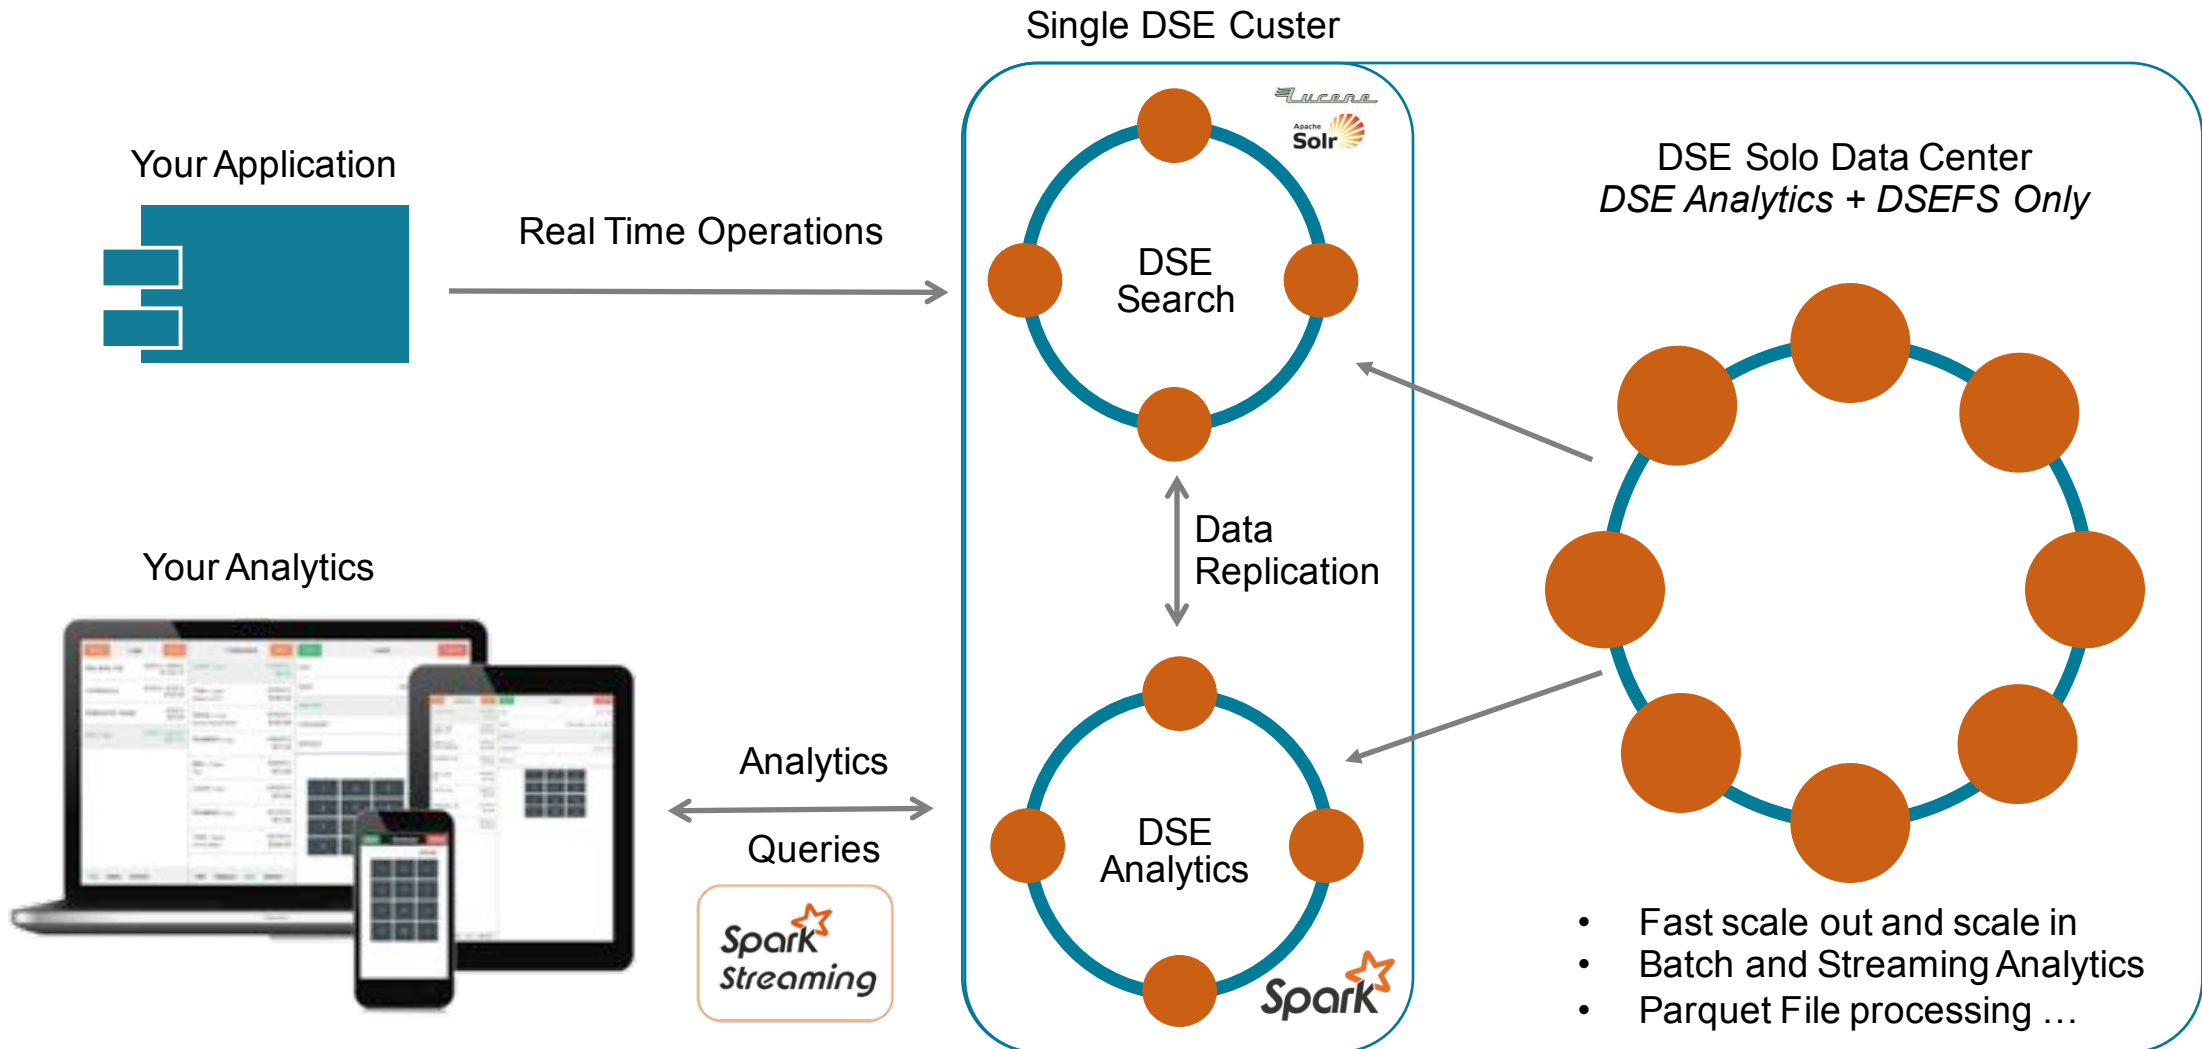
Streaming, ad-hoc, and batch

- High-performance
- Workload management
- SQL reporting

Compared to self-managed:

- No ETL
- True HA without Zookeeper

DSE Reference Architecture



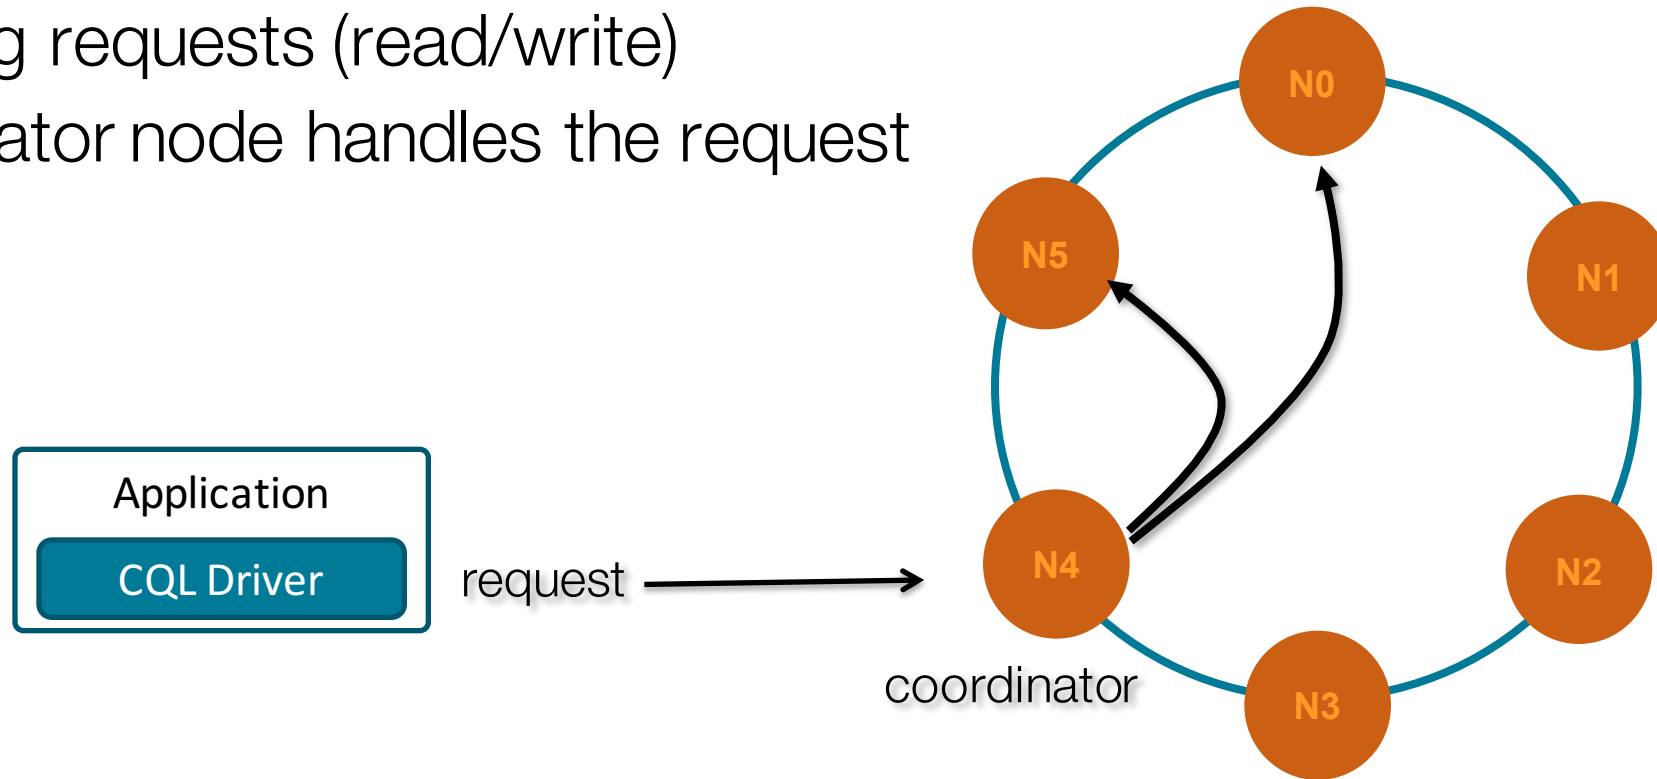
Tunable Consistency

Distributed Read and Write



Data Replication

- Incoming requests (read/write)
- Coordinator node handles the request


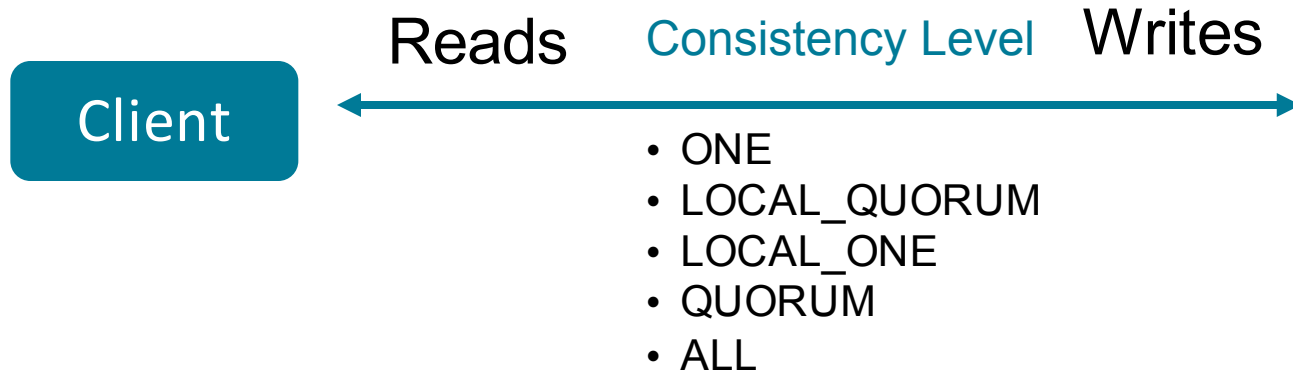


```
CREATE KEYSPACE dsbank WITH replication = {'class': 'NetworkTopologyStrategy', 'DC1': '3'}
```

Every node can be coordinator → masterless

Tunable consistency

- Choose between strong and eventual consistency depending on the need
- Can be done on a per-operation basis, and for both reads and writes
- Handles multi-data center operations
- Light Weight transaction = ACID like



Anti-Entropy and Consistency

Write time


- Tunable Consistency
- Atomic batches
- Hinted handoff

Read Time

- Consistent reads
- Read Repair

Maintenance Time

- Node repair



Lab 1 : Accessing the cluster

Uniform Data Distribution

Query Based Modelling

Primary data model

Name
Value

Column

Partition	Name	Name	Name	...
	Value	Value	Value	...

Row

- Row-oriented, column structure
- Table: similar to an RDBMS table but more flexible/dynamic
 - A row in a table is indexed by its key

Partition	Wide Row Column		Wide Row Column		
	Name	Name	Name	Name	Name	Name	Name	Name	Name	Name	Name	Name
	Value	Value	Value	Value	Value	Value	Value	Value	Value	Value	Value	Value

Wide Row

Modelling explicit partitioning

```
CREATE TABLE messaging.notifications(  
    target_user text,  
    notification_id timeuuid,  
    notification_time timestamp,  
    ...  
    activity text,  
    PRIMARY KEY (target_user, notification_time)  
) WITH CLUSTERING ORDER BY (notification_id DESC)
```

Partition Key (points to `target_user`)

Clustering Column (points to `notification_id`)

Query Optimized (points to `notification_id DESC`)

```
SELECT * FROM notification where target_user = 'mike' limit 1;
```

```
SELECT * FROM notification where target_user = 'mike' AND notification_time >= 2017-11-01 10:00;
```

Primary Key – Unique Identifier

```
PRIMARY KEY ((account_number), transaction_time)
```

Partition Key

- Required to satisfy a queries' predicate(s)
- Ensures row uniqueness
- Defines the location of the partition in the cluster
 - Hashed to ensure even data distribution
- Can be composed of multiple columns
 - “Composite / Compound Key”

Clustering Key

- Sorts data within each partition
 - Defaults to ascending order
- Can Be composed of multiple columns

Modelling explicit partitioning

target_user	notification_id	notification_time	activity
nick	5321998c	2017-11-01 10:00	tom liked
nick	ea1c5d35	2017-11-02 11:00	jake commented
nick	321998c	2017-11-03 09:00	mike created account
mike	e1bd2bcb	2017-11-01 07:00	tom created account

```
SELECT * FROM notification where
target_user = 'nick' AND
target_user = 'mike' AND
notification_time >= 2017-11-01 07:00;
```

Sorted by
notification_time

nick	notification_time: 2017-11-01 10:00		notification_time: 2017-11-02 11:00		notification_time: 2017-11-03 09:00		notification_time: 2017-12-31 23:00	
	ntfcid: 5321998c	activity: tom liked	ntfcid: ea1c5d35	activity: jake commented	ntfcid: 5321998c	activity: mike created account			

Wide Row

Merged, Sorted and Stored Sequentially

Number cells per partition limited to 2 billion
Column max 2 GB , better 1 MB
Blob size 2 GB less then 1 MB recommended



Tokens

Token Range : - 2^{63} to 2^{63}

Data is partitioned after its partition key

A unique token is allocated to a partition


Token = *random* hash of #partition (murmer3)



Data Distribution

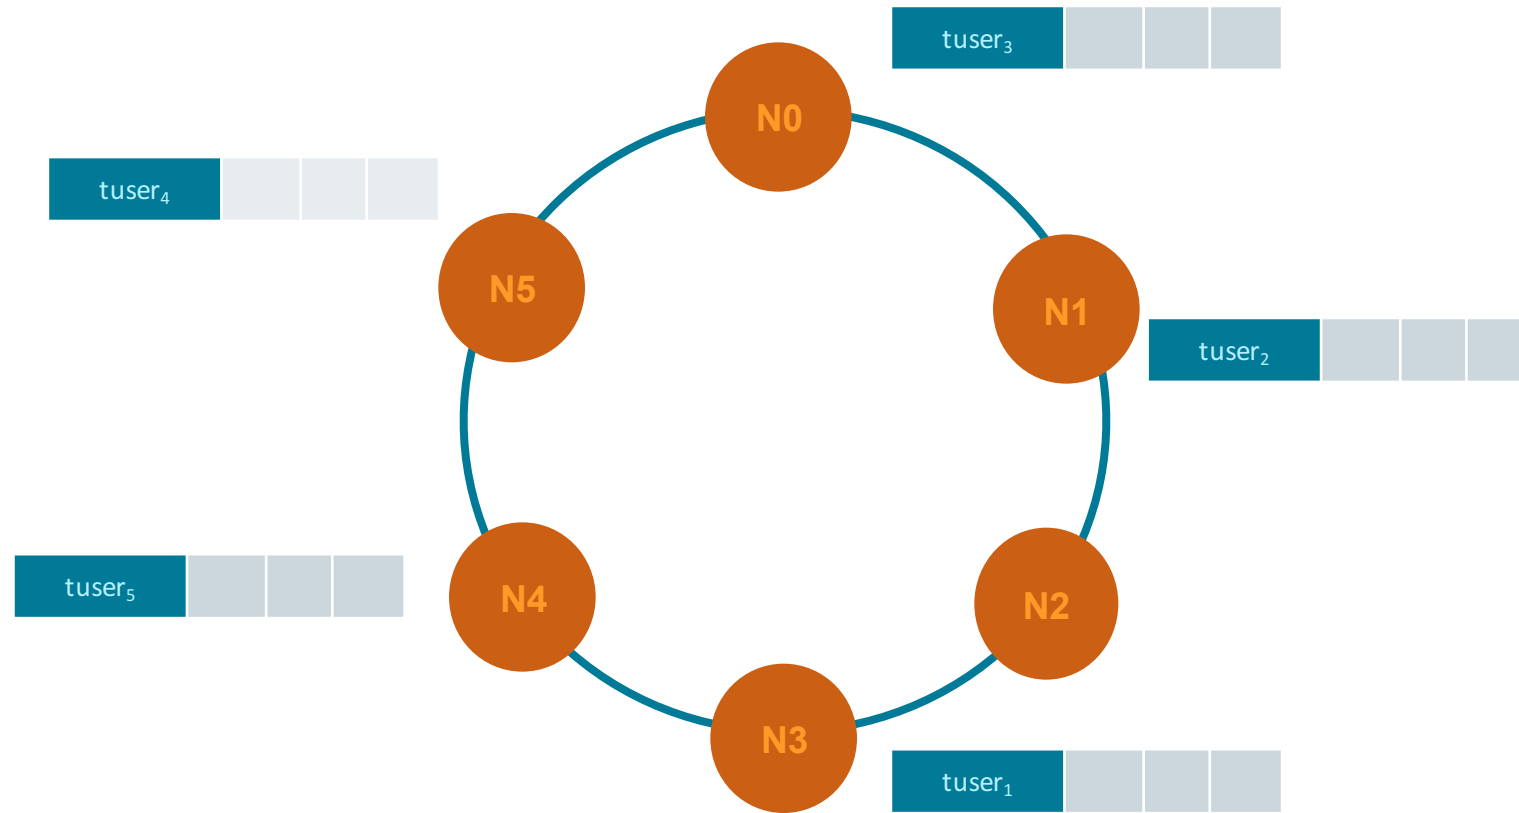
Token = hash of #partition \rightarrow #node

Token1	target_user 1			
Token2	target_user 2			
Token3	target_user 3			
Token4	target_user 4			
Token4	target_user 5			



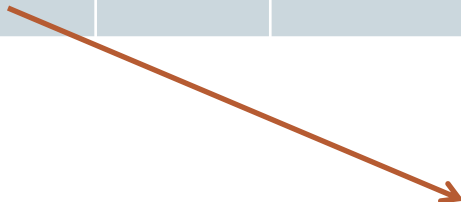
Data is evenly distributed and clock wise replicated

Automated Data Distribution Sharding



What's Stored With Each Column?

nick	notification_time: 2017-11-01 10:00		notification_time: 2017-11-02 11:00		notification_time: 2017-11-03 09:00	
	ntfcd: 5321998c	activity: tom liked	ntfcd: ea1c5d35	activity: jake commented	ntfcd: 5321998c	activity: mike created account



column name : “activity”
column value : “tom liked”
timestamp : 1353890782373000
TTL : 3600

- Last Write Win, cross cluster clock sync, e.g NTP

Skinny vs. Wide rows

Compound Partition Key

PRIMARY KEY ((*target_user*, *day*), notification_time)

- Number cells per partition (2 billion max.)
- Faster operations and lower latency
- multiple gets per dataset if needed
- Equality select

SELECT * FROM notification where target_user = 'mike' AND day IN (1,2,3);



Multiple Clustering Columns

PRIMARY KEY ((target_user), *day*, *notification_time*)

- Simulates 1-N relation ship
- wide rows
- Range selects

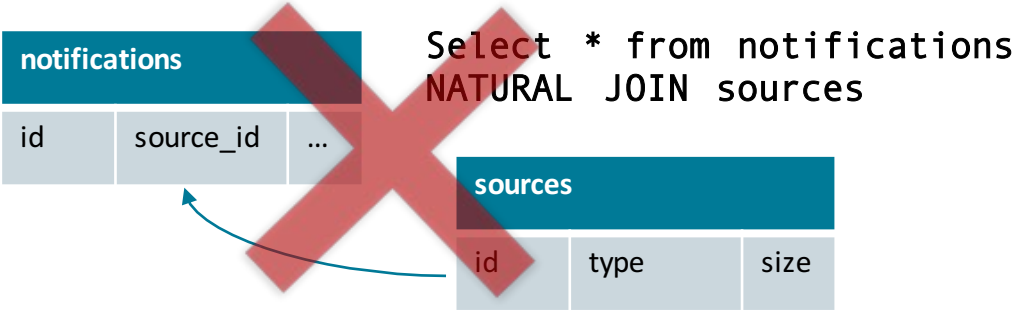
SELECT * FROM notification where target_user = 'mike' AND day IN (1,2,3);

Basic Approach to Data Modeling.

1. What **queries** are needed in your app?
2. What are your **natural unique keys**?
3. Is there **ordering of the data** needed to serve each query?
4. What are the **groupings (1:M, M:M)** in the data?
5. What **filtering** will your queries need?
6. Can events be stored in **chronological order**?
7. Does the **data expire**? Do large **chunks of data expire together**?

Alternatives to joins

- Collections
- Nested frozen Collections
- User Defined Types
- Nested Collections with UDTs
- JSON notation



Collections	User Defined Types			
<Values>	Name	Name	Name	Name
	Value	Value	Value	Value

Partition	Name	Name	Collection	User Defined Types		
	Value	Value	<Values>	Name	Name	Collections
				Value	Value	<Values>

Row

Alternatives to joins

```
CREATE TYPE source_type ( encoding text, size int, location text);
```

```
CREATE TABLE notifications(  
    target_user text,  
    notification_id timeuuid,  
    notification_time timestamp,  
    source map <text, frozen <source_type>>,  
    activity text,  
    PRIMARY KEY (target_user, notification_time)  
) WITH CLUSTERING ORDER BY (notification_time DESC)
```

```
INSERT INTO notifications JSON '{  
    "target_user ": "nick",  
    "notification_id ": "5321998c",  
    "notification_time ": "2017-11-01",  
    "source": {  
        "profile_pic": {  
            "encoding": "jpeg",  
            "size": 15,  
            "location": "/"}}}}';
```

DSE Performance Basics

Can DSE be both bigger and faster? Yes it can.

More Throughput? More Data?

Use more nodes (scale out)

Do not use too big nodes (scale ~~up~~)

Know Cassandra ops best practices
Use OpsCenter to monitor, alert, repair

Faster Operations? Predictable Latency?

Check your data model and queries

Use asynchronous queries

Use prepared statements

Compaction tuning or maybe strategy

Lab 2 : DSE Core and Operations

Questions?

Negib Marhoul
Solutions Engineer
negib.marhoul@datastax.com

Thank you

Günther Schnack
Enterprise Sales Manager
PLZ: 7 – 9 & 5
+44 79 49 448 969
gunther.schnack@datastax.com

Jerome Ruiz
Regional Channel Director
jerome.ruiz@datastax.com

Negib Marhoul
Solutions Engineer
negib.marhoul@datastax.com

Thank you

Giscard Venn
Enterprise Sales Manager
PLZ: 0 - 4 & 6
+49 173 4516618
giscard.venn@datastax.com

Jerome Ruiz
Regional Channel Director
jerome.ruiz@datastax.com